

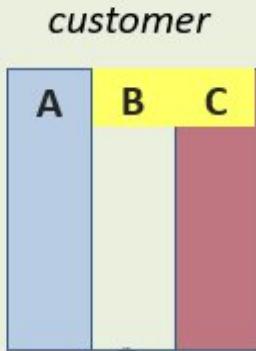


Vertica

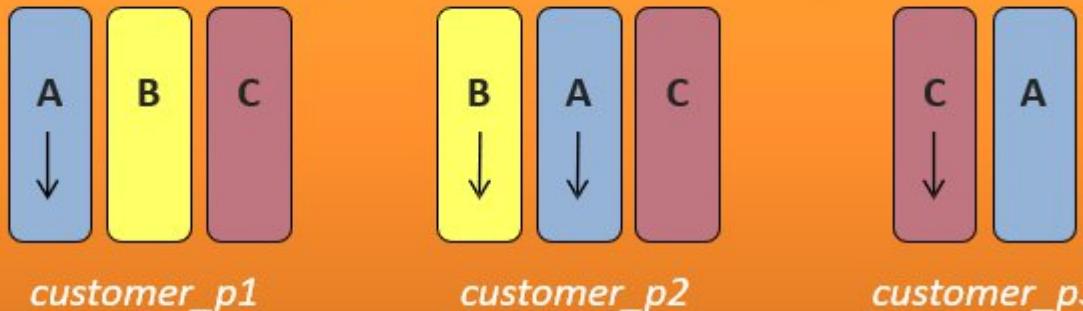
Logical and Physical Representations

Logical and Physical Representations

TABLE
Logical



PROJECTIONS
Physical



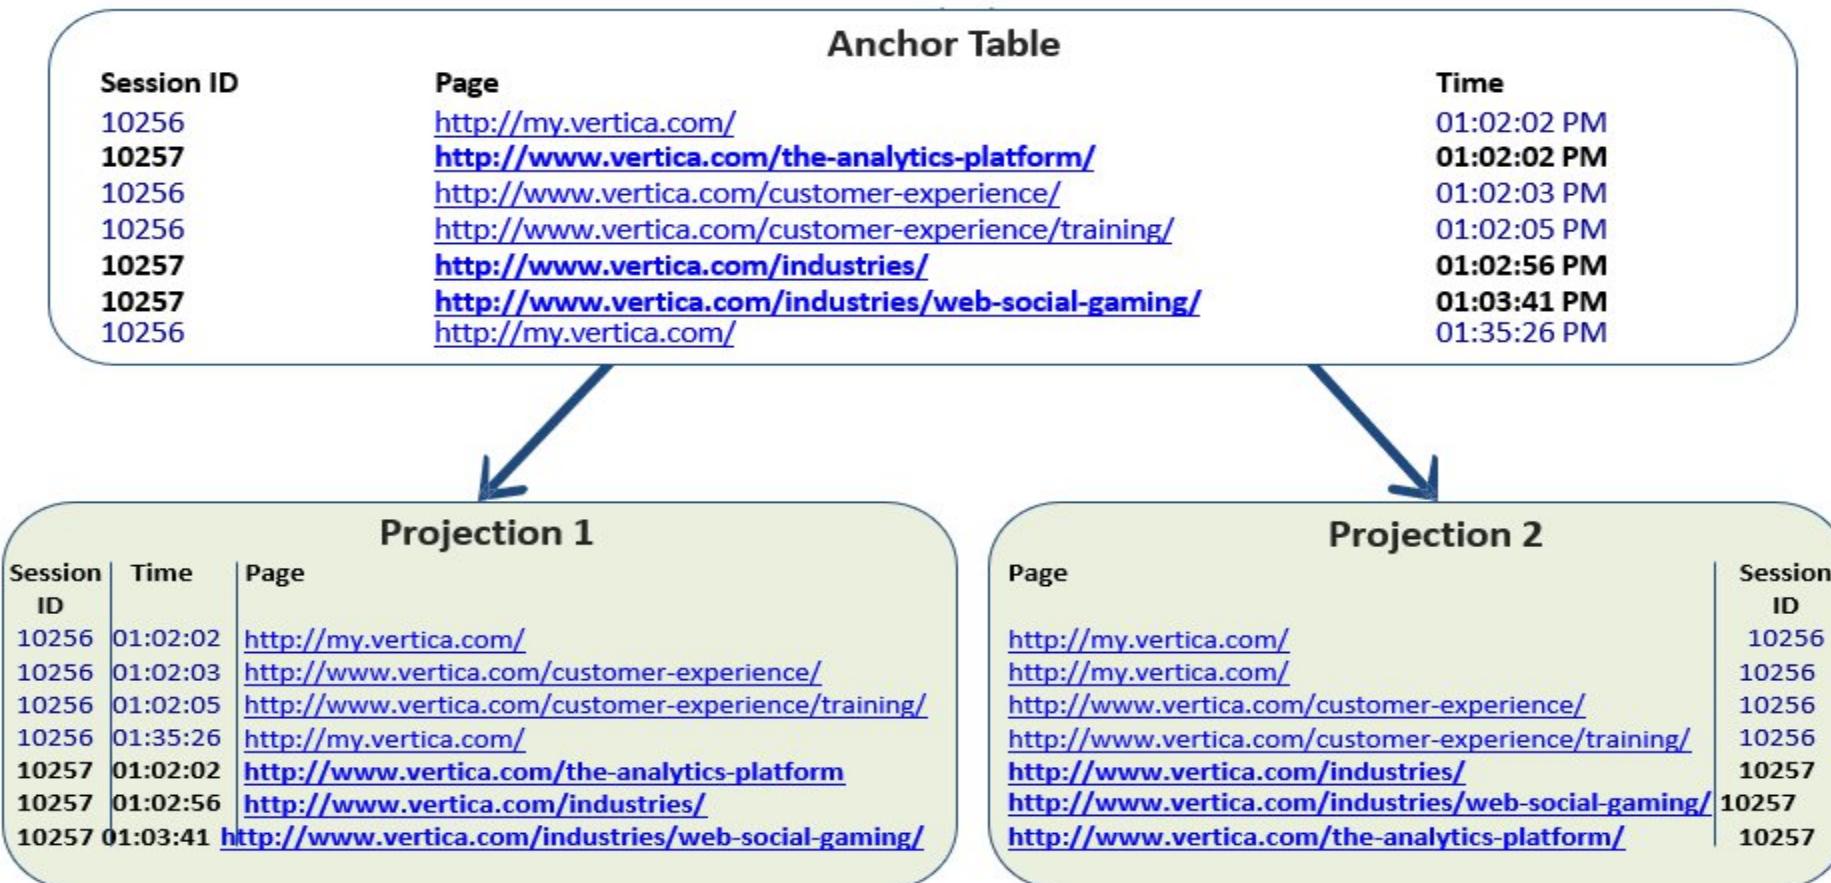
Tables

- Tables are logical objects in Vertica
- Queries reference tables
- Optimizer picks best projection(s) to answer queries
- Tables are owned by a user and created in a schema
- \d to see all tables

Types of Temporary Tables

- GLOBAL – default, DDL visible in all sessions
- LOCAL – DDL visible only in the session created

Tables versus Projections



Projections are optimized for common query patterns

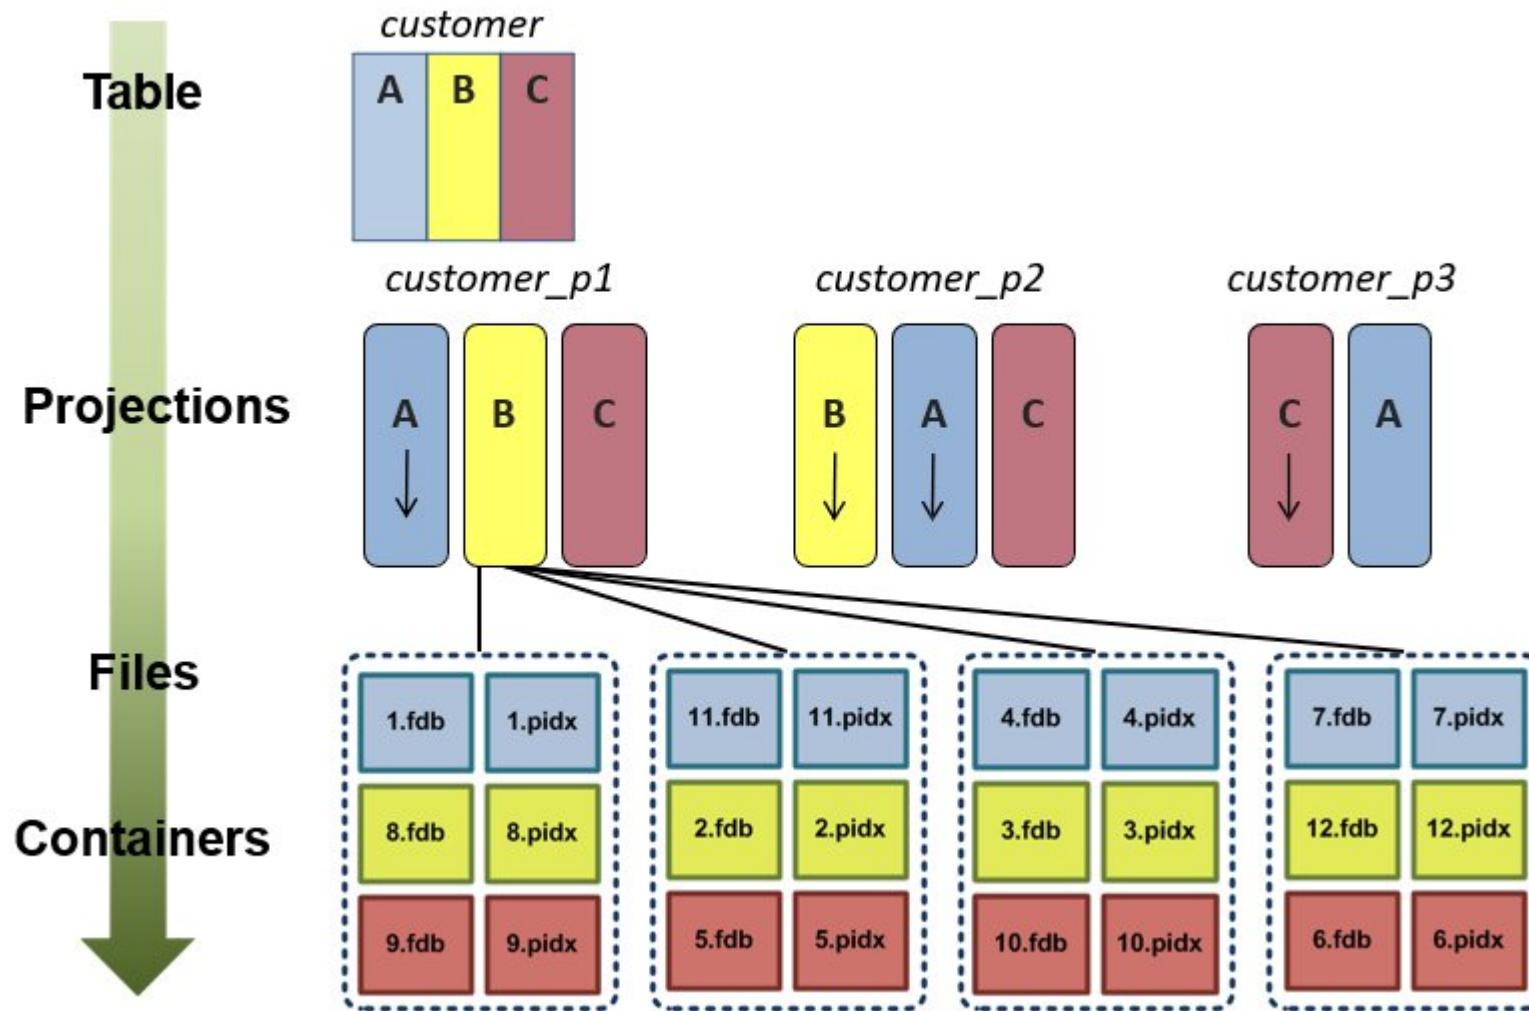
What is a Projection?

- Collection of table columns
- Stores data in a format to optimize query execution
- Similar in concept to materialized views

Projection Basics

- Anchor table is not stored
- Data is stored in a sorted and compressed format
- No need for indexing
- Transparent to the end user and applications
- Created by Database Designer (DBD)
 - Can be manually created
- Best projection for a query is chosen by the Optimizer at query execution time

Vertica Object Hierarchy



Projection Types

- Superprojection
 - Projection that contains all columns of a logical table
- Query-specific projection
 - Optimized for a specific query or class of queries
 - Can include a subset of columns
- Buddy projection
 - Projection with the same columns and segmentation on different nodes to provide high availability (HA)

How to Create Projections

- Automatically on first data load into a table
 - Unoptimized superprojections
- Automatically using the Database Designer
 - Optimized superprojections
 - Possible query-specific projections
- Manually using CREATE PROJECTION syntax in a vsql statement

Projection Properties

- Encoding/Compression
- Sort Order
- Replication or Segmentation
- K-safety

Encoding/Compression and Sorting

- Each column is always encoded, compressed, or both
- Vertica can work directly with encoded data; compressed data must first be uncompressed
- All projections contain at least one column in the ORDER BY statement

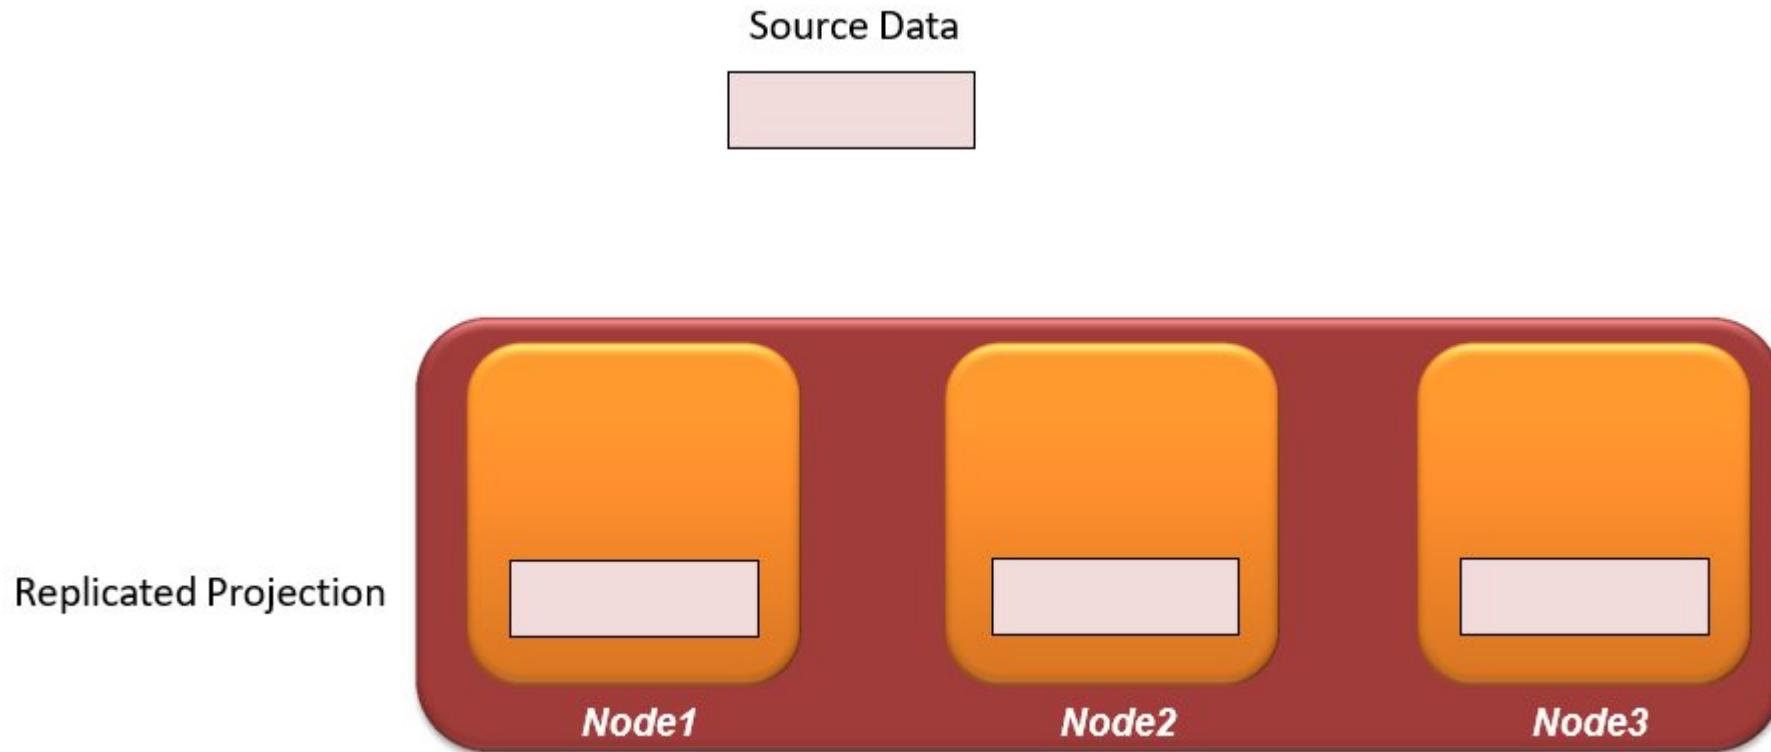
ORDER BY A		
A	B	C
1	2	b
1	1	a
1	3	c
2	1	d
2	3	f
2	2	e

ORDER BY AB		
A	B	C
1	1	a
1	2	b
1	3	c
2	1	d
2	2	e
2	3	f

ORDER BY B		
A	B	C
1	1	a
2	1	d
3	2	b
4	2	e
2	3	f
1	3	c

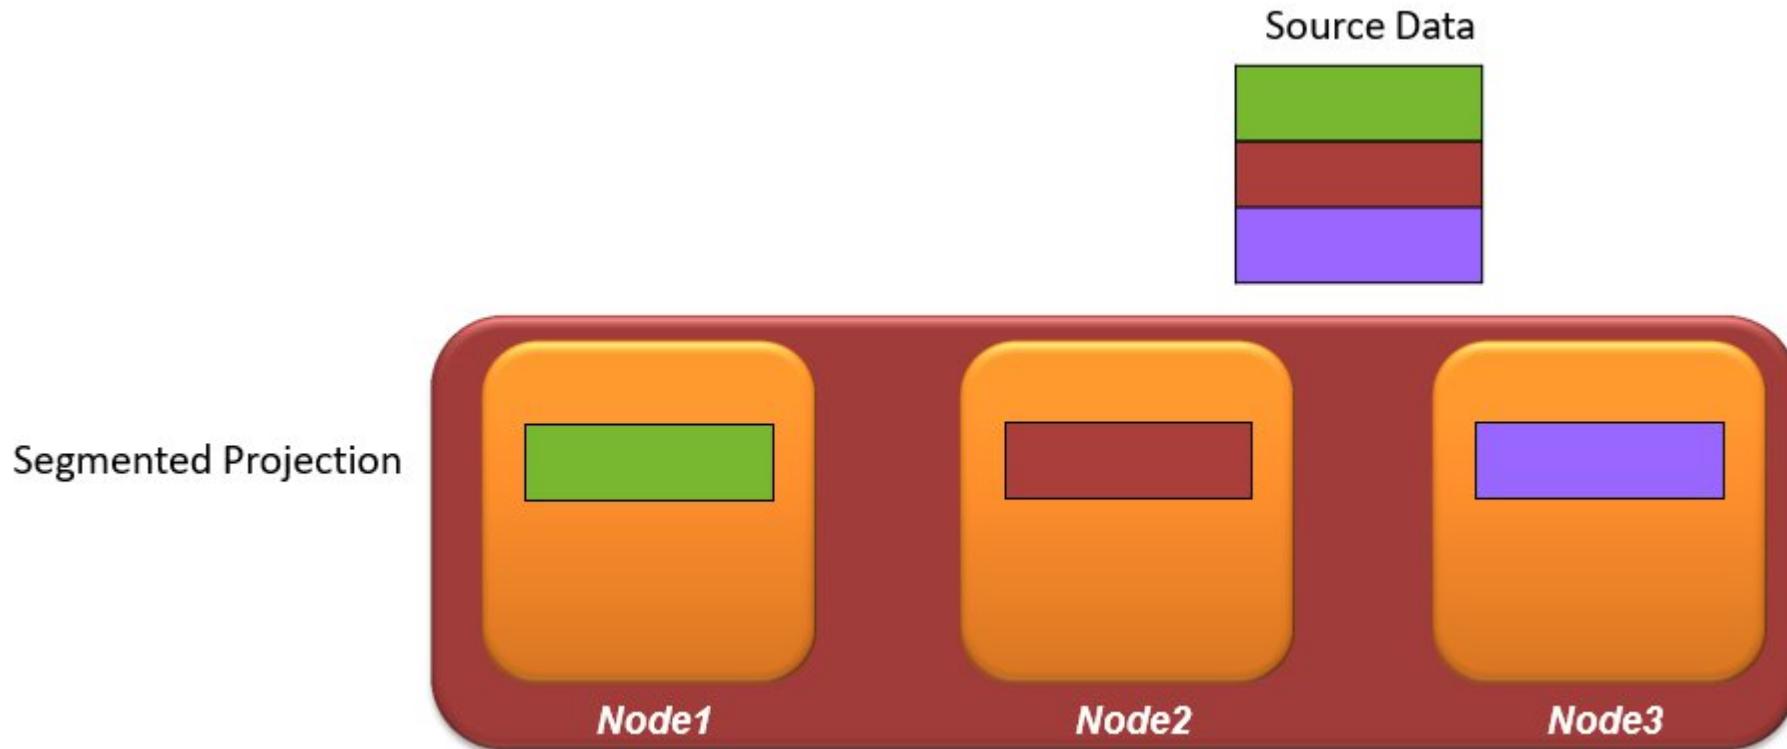
Replication

- For a small projection, copy the full projection to each node
 - Inherently provides high availability of the projection



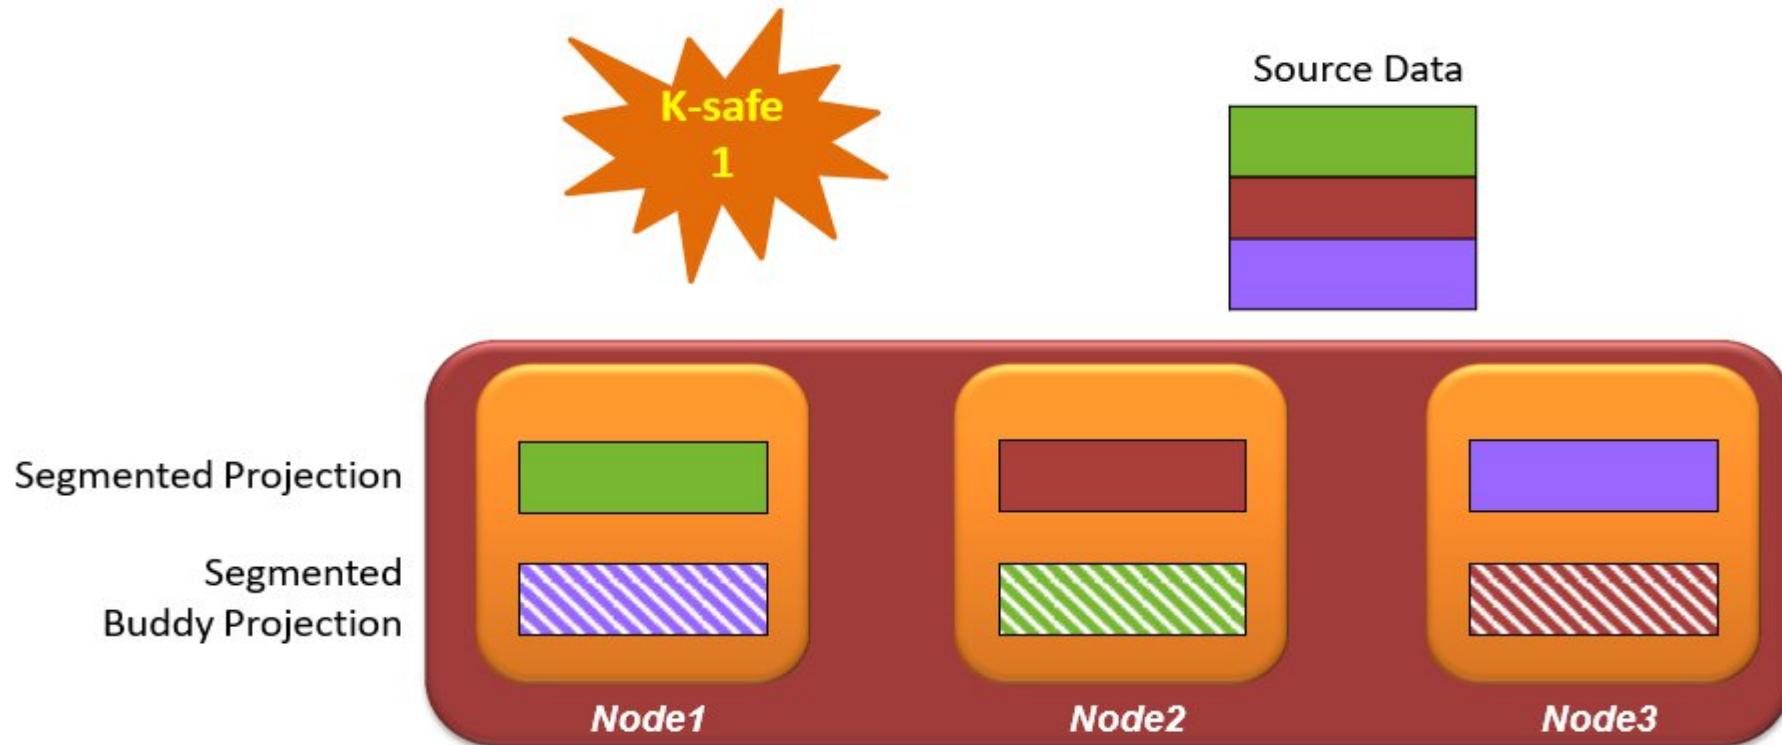
Segmentation

- For a large projection, distribute projection data across multiple nodes
 - Segmenting the projection splits the data into **subsets** of rows
 - Design for **random** data distribution



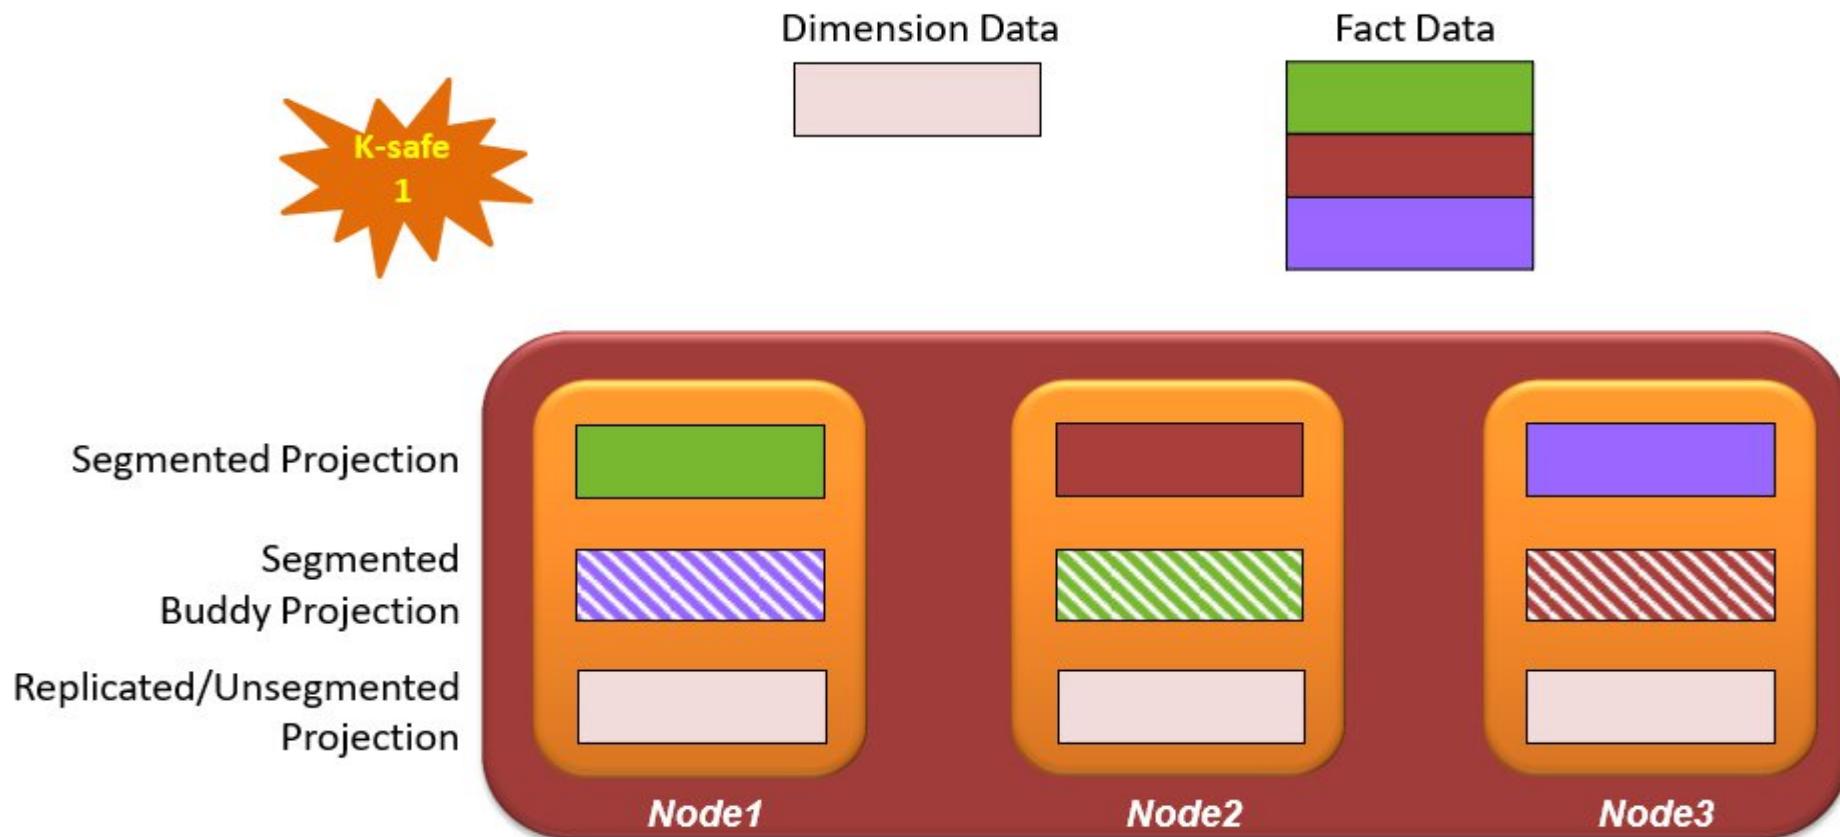
Segmentation and High Availability

- Buddy projections provide high availability for segmented projections
 - Buddy projection is a copy of a projection segment that is stored on a different node
 - K-safety is the number of replicas of data that exist on the cluster



Segmentation, Replication, K-safety

- Segmentation with buddy projections
- Replication



Create Projection DDL

```
CREATE projection snmp_p1 (
    host      encoding rle,
    interface encoding rle,
    time      encoding commanddelta_comp,
    metric    encoding rle,
    value     encoding deltaval
) as
```

Column List
and Encodings

```
SELECT
    host,
    interface,
    time,
    metric,
    value
FROM snmp
```

Base Query (Joins allowed)

```
ORDER BY host, interface, metric, time
```

Sort Order

```
SEGMENTED BY hash ( host, interface,
    time, metric ) ALL NODES KSAFE;
```

Segmentation
and K-safety

Projection Basics: Maintenance

- Data is loaded directly into projections
- No need to rebuild or refresh projections after the initial refresh
- New projections can be created at any time, either manually or by running the Database Designer
- Old projections can be dropped

The background features a series of glowing blue light streaks and curves, resembling motion blur from a camera or light painting, set against a dark, solid black background.

opentext™

Thank you