



LOADING DATA INTO VERTICA

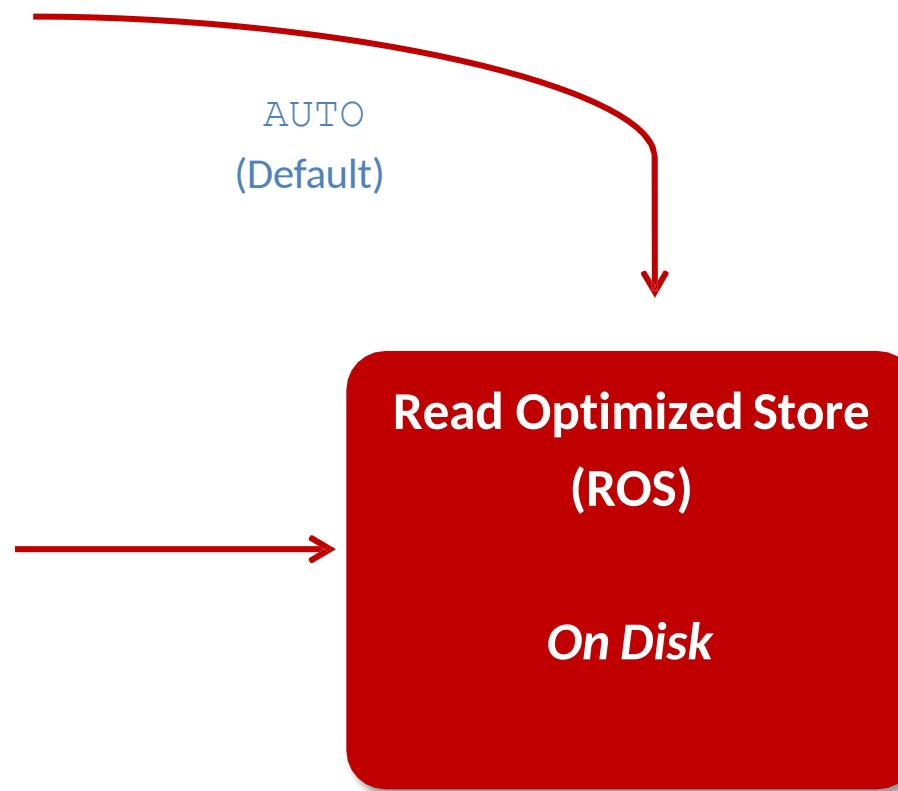
# Loading to ROS: Trickle Loading and Bulk Loading

- Trickle Load

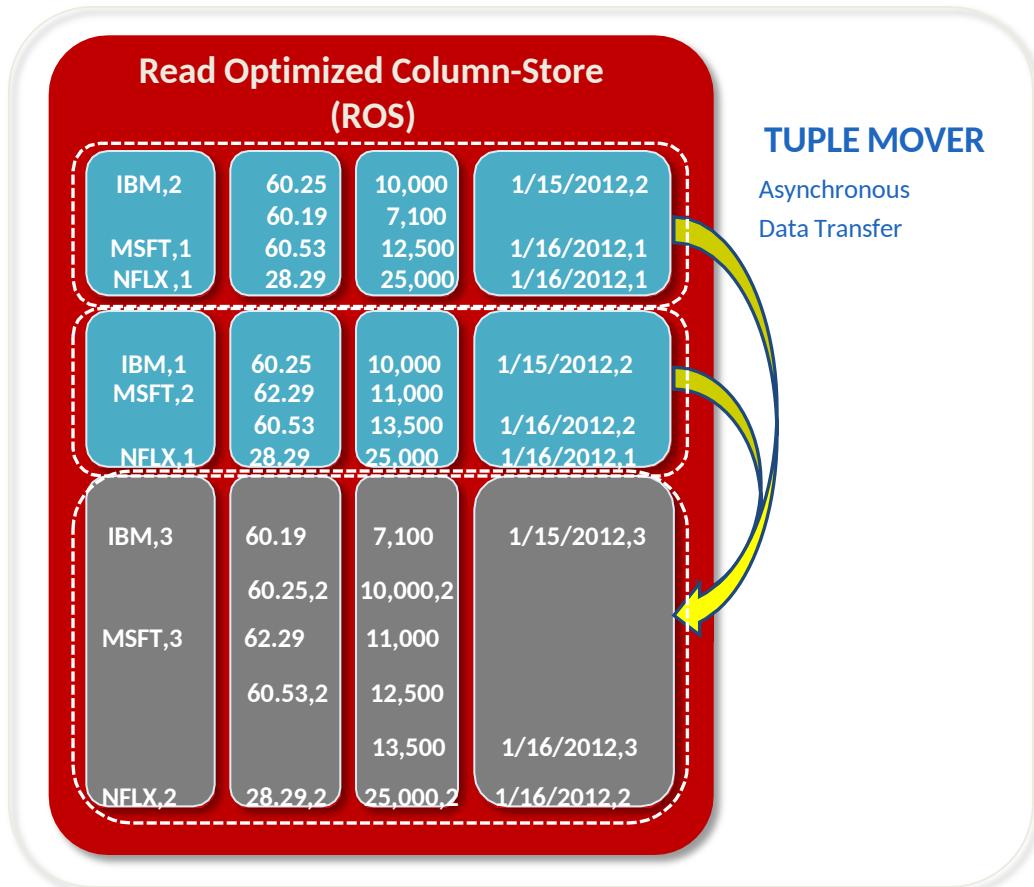
- INSERT
- UPDATE
- DELETE
- COPY

- Bulk Load

- INSERT /\*+DIRECT\*/
- UPDATE /\*+DIRECT\*/
- DELETE /\*+DIRECT\*/
- COPY DIRECT



# Tuple Mover - mergeout



- ROS created by COPY
- Defragments storage containers by merging them
- Purges deleted records
- Honors projection definition
- Maintains partition boundary on active partitions

# Tuple Mover – mergeout parameters

- ActivePartitionCount
  - default 1, increase if simultaneously loading multiple partitions
- MergeOutInterval
  - default 10 minutes, periodic check for mergeout

The background features a series of glowing blue light streaks and curves against a black background, creating a sense of motion and depth.

# opentext™

Vertica : Adding Rows

# Adding rows: INSERT, COPY

- INSERT
  - Use infrequently, high overhead
  - If necessary `INSERT /*+DIRECT*/ ... SELECT ...`
- COPY
  - Faster approach to load a warehouse
  - Can load in parallel

# COPY

- Bulk loads data from files to Vertica
  - COPY FROM STDIN pipes from STDIN
    - zcat table\_a.gzip | vsql -c "COPY table\_a from STDIN DIRECT;"
  - COPY FROM file loads from Vertica nodes
- Usually TEXT loading, NATIVE (via drivers)
- See SQL Reference manual for command options

# COPY: Checking Data Format Before Loading

- \$ file Date\_Dimension.tbl  
Date\_Dimension.tbl: ASCII text
- \$ wc Date\_Dimension.tbl  
1828 5484 221822 Date\_Dimension.tbl
- \$ file data\*  
data1.txt: Little-endian UTF-16 Unicode text  
data2.txt: ISO-8859 text

# COPY: Converting Files Before Loading Data

- `iconv -f ISO88599 -t utf-8 data2.txt > data2-utf8.txt`

# COPY: Checking UTF-8 Compliance After Loading Data

- SELECT name FROM nametable WHERE ISUTF8(name) = FALSE;

# COPY: Loading Data Interactively

- \$ cat fact\_table.tbl | vsql -c "COPY FACT\_TABLE FROM STDIN DELIMITER '|' DIRECT";
- \$ cat fact\_table.tbl | vsql -c "COPY FACT\_TABLE FROM LOCAL STDIN DELIMITER '|' DIRECT";

# COPY LOCAL

- Enables copying of a file stored on a local machine
  - dbadmin access NOT required
  - Exceptions and rejections can be directed to the client machine
  - No need to set up permissions on Storage Locations
- COPY <file> from LOCAL '/home/user/data/\*'
- => COPY store.store\_dimension FROM LOCAL '/usr/files/my\_data/input\_file' GZIP;
- =>COPY simple\_table FROM LOCAL 'input\_file.bz' BZIP,  
'input\_file.bz' BZIP;
- => COPY simple\_table FROM LOCAL STDIN;

# COPY NO COMMIT

- COPY automatically commits by default
  - COPY NO COMMIT to prevent this
  - Enables review of error logs before deciding to commit
  - Combine Multiple COPY Statements in Same Transaction

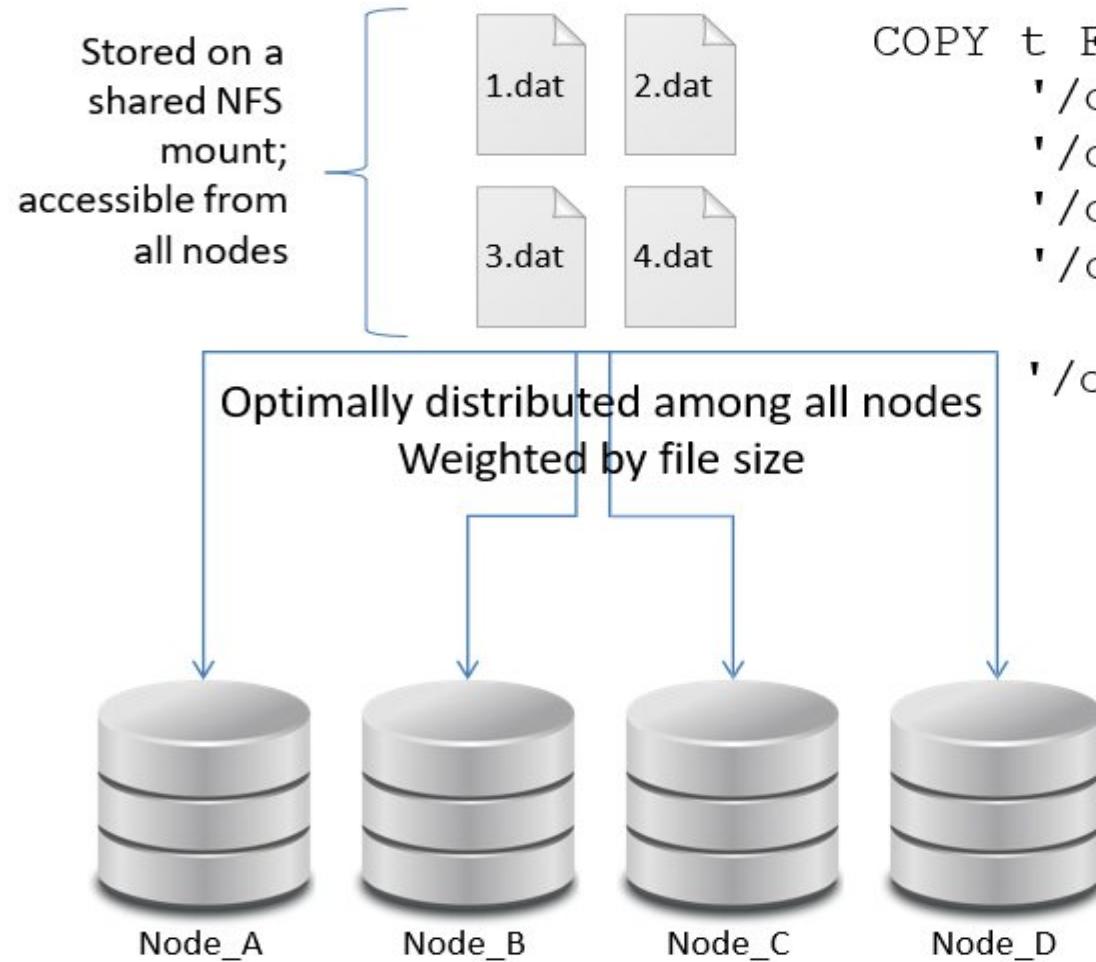
```
COPY... NO COMMIT;  
COPY... NO COMMIT;  
COPY... NO COMMIT;  
COPY X FROM LOCAL NO COMMIT;  
COMMIT;
```

# COPY Command Options

- Columns as expressions
  - ```
COPY Retail.Dim (num, name, store, date as SYSDATE)
FROM '/home/dbadmin/dim.txt' DELIMITER '|';
```
- Columns computed from other columns
  - ```
COPY Retail.Dim (num, name, store,
joined filler date,
joined_mth AS DATE_PART('month',joined), joined_year AS
DATE_PART('year',joined)) FROM '/home/dbadmin/dim3.txt'
DELIMITER '|';
```
- Expressions CAN use most SQL functions, operators, constants, NULLs, and comments
- Expressions CANNOT use Vertica analytic or aggregate functions

# Multi-file COPY

## COPY ... ON ANY NODE



```
COPY t FROM
  '/data/1.dat' ON ANY NODE,
  '/data/2.dat' ON ANY NODE,
  '/data/3.dat' ON ANY NODE,
  '/data/4.dat' ON ANY NODE;
OR
'/data/*.dat' ON ANY NODE;
```

## Multi-file COPY

- ▶ COPY t FROM '/data/file1.dat' ON v\_vmart\_node0001, '/data/file2.dat' ON v\_vmart\_node0002;
- ▶ COPY t FROM '/data/\*.dat' ON ANY NODE;
- ▶ COPY t FROM '/data/bigfile.dat' ON ANY NODE;
- ▶ => COPY t FROM '/data/big1.dat' ON (v\_vmart\_node0001, v\_vmart\_node0002, v\_vmart\_node0003), '/data/big2.dat' ON (v\_vmart\_node0004, v\_vmart\_node0005);

The background features a series of glowing blue light streaks and curves against a black background, creating a sense of motion and depth.

# opentext™

Vertica : Updating Rows

# Updating Rows: Update/Delete

- Delete
  - Check performance
  - Not performed in place
  - PURGE removes files
- Update
  - More efficient to split into delete and bulk load
  - Use delete best practices.

# Updating Rows: Bulk Delete

- Extract all delete into a file
- => CREATE LOCAL TEMP TABLE data\_to\_delete (emp\_id INT);
- => COPY data\_to\_delete FROM '/tmp/employee\_to\_delete.txt' ;
- => DELETE /\*+ direct \*/ FROM store.store\_orders\_fact WHERE employee\_key IN (SELECT \* FROM data\_to\_delete );
- => DROP TABLE data\_to\_delete ;

## Updating Rows: PURGE

- Permanently removes delete vectors from ROS storage containers so disk space can be reused. PURGE removes all historical data up to and including the Ancient History Mark epoch.
- SELECT PURGE();
- SELECT PURGE\_TABLE('store.store\_sales\_fact');

# Bulk Update Recommendation

1. Store new batch by loading the new batch in staging table
2. Identify update records by joining batch with main table and extract keys for obsolete records
3. Store keys in separate table
4. Delete records through the support of a subquery
5. Bulk load the staging table
6. Clean up temporary tables

# MERGE

- Data being loaded includes both new and existing data
- Inserts new records, updates existing records

New Data (source)

UserID	X	Y	Count	Name
1	10.1	2.7	1	AAA
2	5.1	7.9	1	BBB
3	4.1	7.7	1	CCC

Existing Data (target)

UserID	X	Y	Count	Name
1	10.1	2.7	1	AAA
1	4.1	7.7	1	CCC
2	4.1	7.7	1	CCC

MERGE  
into

Result

UserID	X	Y	Count	Name
1	10.1	2.7	2	AAA
1	4.1	7.7	1	CCC
2	4.1	7.7	1	CCC
2	5.1	7.9	1	BBB
3	4.1	7.7	1	CCC

# MERGE Syntax

```
MERGE INTO location target USING new location source  
ON source.userid = target.userid AND source.x = target.x  
AND source.y = target.y WHEN MATCHED THEN  
UPDATE SET count = target.count + source.count  
WHEN NOT MATCHED THEN  
INSERT VALUES (src.userid, src.x, src.y, src.count, src.name);
```

The background features a series of glowing blue light streaks and curves, creating a sense of motion and depth against a dark background.

# opentext™

Vertica : Post Loading Tasks

# Update Statistics

- Data statistics written to catalog for use by query Optimizer
- When to run?
  - After an initial data load
  - If data has changed > 50%
  - If query performance changes over time (stale statistics)

# Update Statistics

- How to run?

```
SELECT analyze_statistics ('table-name');
```

- Use " for all tables

- How to check statistics?

```
SELECT get_projection_status ('projection-name');
```

```
SELECT has_statistics FROM projections;
```

- Specific recommendations appear in Workload Analyzer in Management Console

The background features a series of glowing blue light streaks and curves, resembling motion blur from a camera or light painting, set against a dark, solid black background.

# opentext™

Thank you