

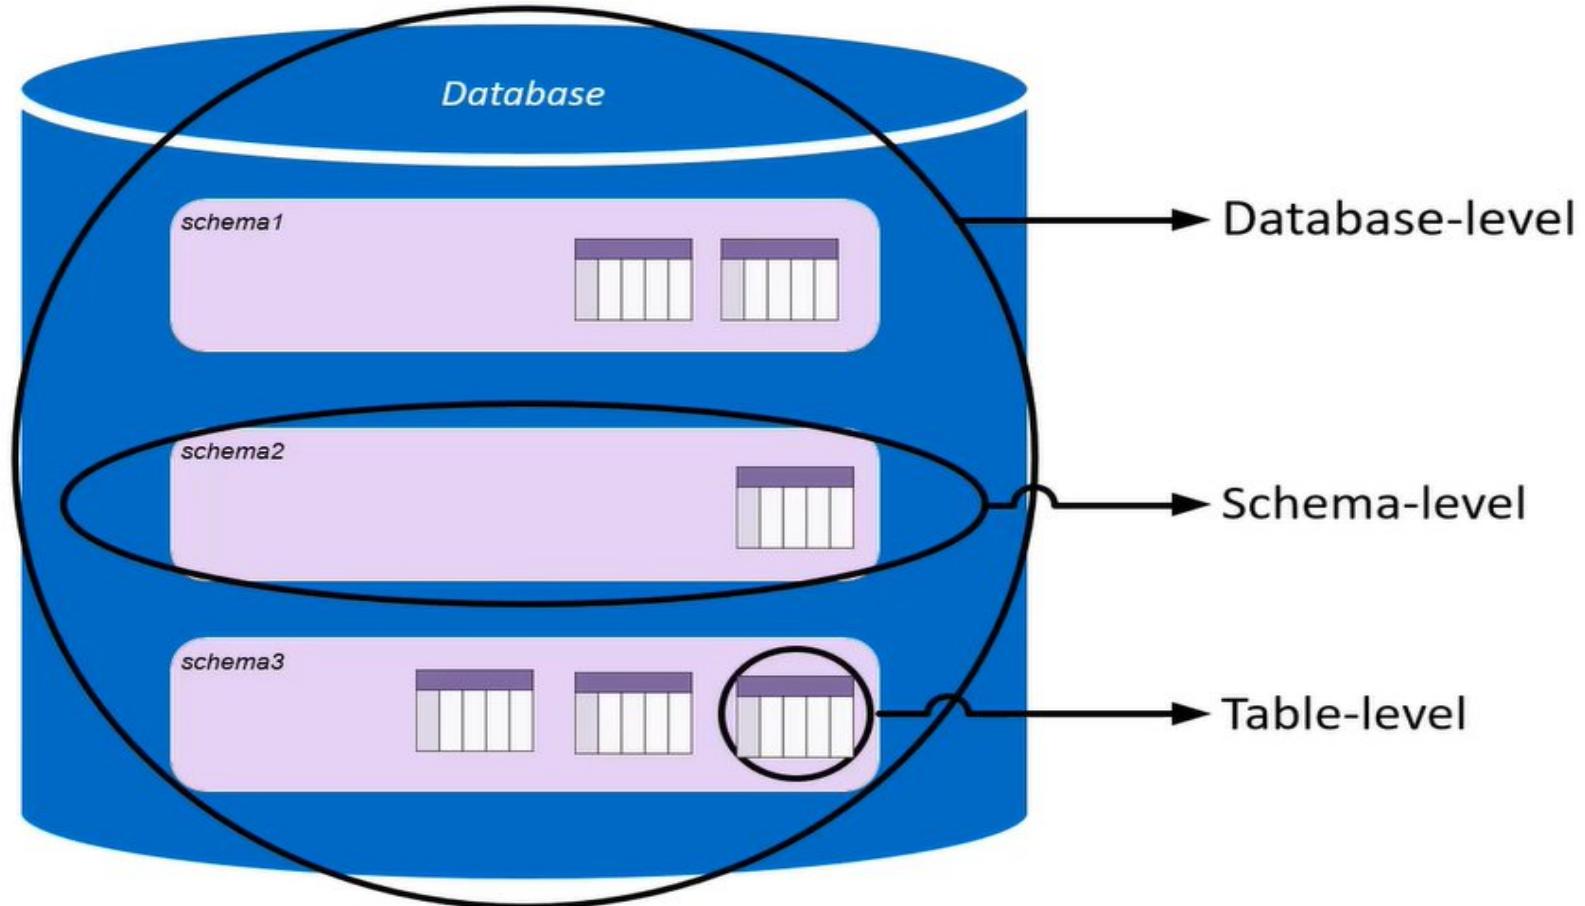


Backup and Restore

Levels of Data Protection

Backup Strategy	Advantages/Disadvantages	Resource Cost
Back up to external storage site <ul style="list-style-type: none">• Full backup with full/object-level restore to current cluster• Object-level backup/restore to current cluster	<ul style="list-style-type: none">• Protects against user error, logical corruption, application error, hardware failure• Full or granular incremental backups• Requires allocation of hardware resources	Medium
Back up to local hard-link copy <ul style="list-style-type: none">• Copy of catalog, hard link to data	<ul style="list-style-type: none">• Protects against user error, logical corruption, application error• Fast and space-efficient• No hardware failure protection	Low
Replication to remote data center <ul style="list-style-type: none">• Full database copy, object-level restore to remote cluster	<ul style="list-style-type: none">• Protects against user error, logical corruption, application error, hardware failure, data center failure• Online replication to active target database• Requires managing another data center	High

Database Backup Levels



Database Backup Levels

- On a regular schedule, as part of regular database maintenance
- Before and after upgrading Vertica
- Before and after a single load of a large volume of data
- Before dropping partitions
- Before adding, removing, or replacing nodes
- After recovering a cluster from a crash

Backup – Overview

- Vertical data files are write-once
- Number of files increases with each load
- Tuple Mover keeps the number of files under control
- To backup, copy Vertical files to stable storage

Backup and Restore Options

- Backup and Restore by Database
 - General backup process
- Backup and Restore by Schema
 - Multi-tenant database with different backup frequency
 - Backup and Restore by Table
 - Backup data specific to Vertica
 - Verify space, restore single table, etc.

Creating Backup Configuration File

- Located at /opt/vertica/bin/vbr.py
- Create configuration file first
 - vbr.py - -setupconfig
 - Defines where the database backup is saved, the temporary directories to use, and which nodes, schema(s), and/or table(s) in the database are to be backed up

Sample Backup Configuration File

backup_restore_full_external.ini

```
[dbadmin@node1 ~]$ more /opt/vertica/share/vbr/example_configs/backup_restore_full_external.ini
; This sample vbr configuration file shows full or object backup and restore to a separate remote backup-host for
each respective database host.
; Section headings are enclosed by square brackets.
; Comments have leading semicolons (;) or pound signs (#).
; An equal sign separates options and values.
; Specify arguments marked '!!!Mandatory!!!' explicitly.
; All commented parameters are set to their default value.

; -----
;::: BASIC PARAMETERS ::;
; -----


[Mapping]
; !!!Mandatory!! This section defines what host and directory will store the backup for each node.
; node_name = backup_host:backup_dir
; In this "parallel backup" configuration, each node backs up to a distinct external host.
; To backup all database nodes to a single external host, use that single hostname/IP address in each entry below.

v_exampledb_node0001 = 10.20.100.156:/home/dbadmin/backups
v_exampledb_node0002 = 10.20.100.157:/home/dbadmin/backups
v_exampledb_node0003 = 10.20.100.158:/home/dbadmin/backups
v_exampledb_node0004 = 10.20.100.159:/home/dbadmin/backups

[Misc]
; !!!Recommended!! Snapshot name. Object and full backups should always have different snapshot names.
; Backups with the same snapshotName form a time sequence limited by restorePointLimit.
; SnapshotName is used for naming archives in the backup directory, and for monitoring and troubleshooting.
; Valid characters: a-z A-Z 0-9
; snapshotName = backup_snapshot

[Database]
; !!!Recommended!! If you have more than one database defined on this Vertica cluster, use this parameter to specify
; which database to backup/restore.
; dbName = current_database

; If this parameter is True, vbr prompts the user for the database password every time.
; If False, specify the location of password config file in 'passwordFile' parameter in [Misc] section.
; dbPromptForPassword = True

; -----
;::: ADVANCED PARAMETERS ::;
; -----


[Misc]
; The temp directory location on all database hosts
```

Backup Preparation (1 of 2)

- Verify your database is running
- All of the backup hosts must be up and available
- The location to store the backups must have sufficient disk space to store the backups, and must be writable by the user account used to start the backup utility

Backup Preparation (2 of 2)

- Backups are stored in the location you specify in the configuration file you use to back up the database
- The directory containing the backup file has a subdirectory for each node backed up to that location, which in turn contains a directory with the name of the backup snapshot
 - The snapshot name is set using the snapshotName option in the configuration file

Initializing the Backup Location

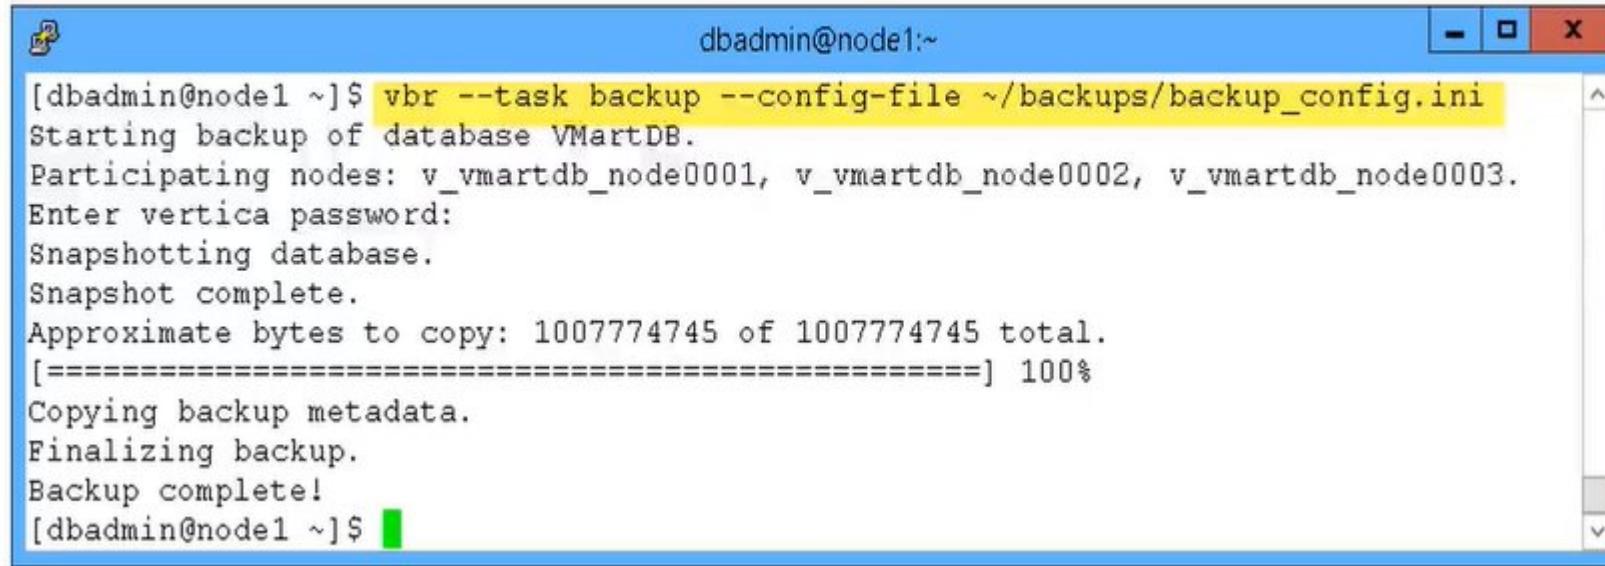
```
dbadmin@node1:~ [dbadmin@node1 ~]$ vbr --task init --config-file ~/backups/backup_config.ini  
Initializing backup locations.  
Backup locations initialized.  
[dbadmin@node1 ~]$
```

```
dbadmin@node1:~ [dbadmin@node1 ~]$ more ~/backups/backup_manifest  
[snapshots]  
[objects]  
[dbadmin@node1 ~]$
```

Performing a Backup

- vbr.py - -task backup - -config-file <configfile>
used for full and incremental backups
- First run does a full backup
- Subsequent runs only copy files added since the last backup
 - Vertica's files are write-once
 - Files are only added or deleted, never modified

Running the Backup



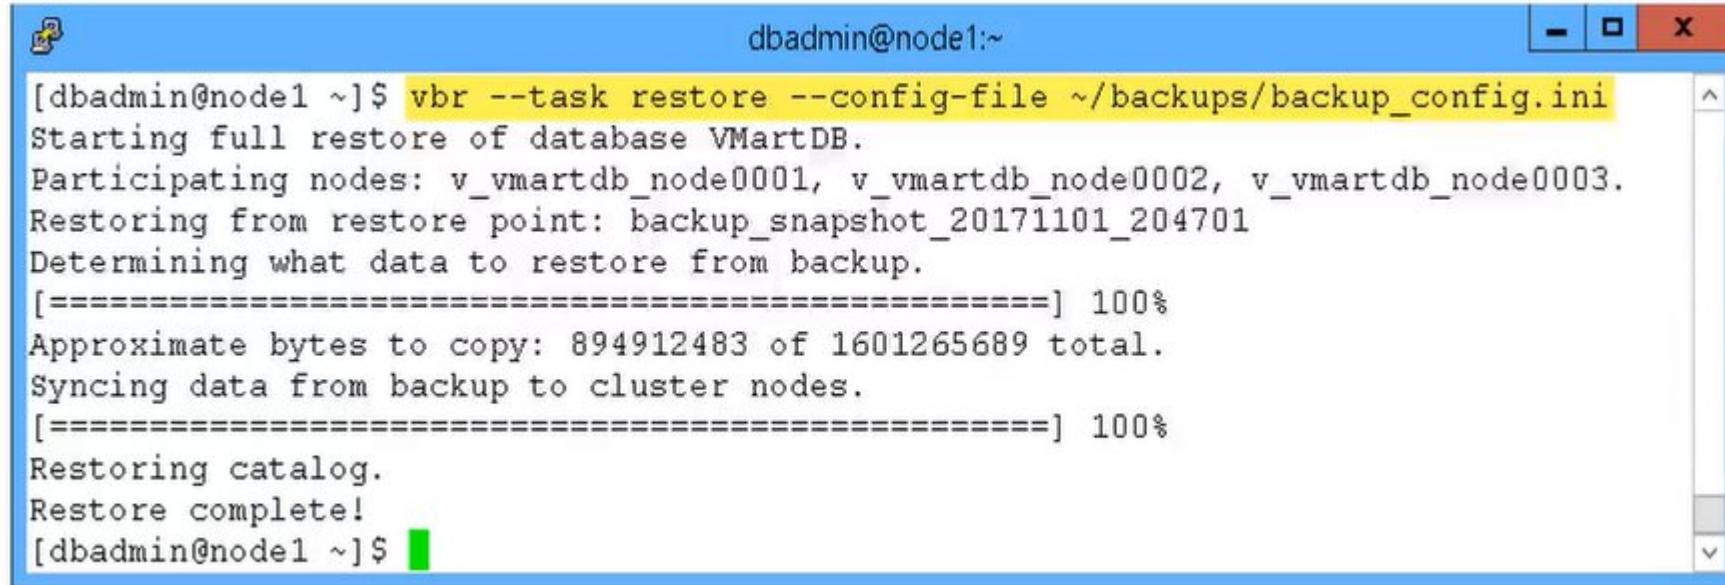
A screenshot of a terminal window titled "dbadmin@node1:~". The window contains the following text:

```
[dbadmin@node1 ~]$ vbr --task backup --config-file ~/backups/backup_config.ini
Starting backup of database VMartDB.
Participating nodes: v_vmartdb_node0001, v_vmartdb_node0002, v_vmartdb_node0003.
Enter vertica password:
Snapshotting database.
Snapshot complete.
Approximate bytes to copy: 1007774745 of 1007774745 total.
[=====] 100%
Copying backup metadata.
Finalizing backup.
Backup complete!
[dbadmin@node1 ~]$
```

Performing a Restore

- vbr.py - -task restore - -config-file <configfile>
- Interactive prompts guide the restore process
- Can restore entire database, schema, or table

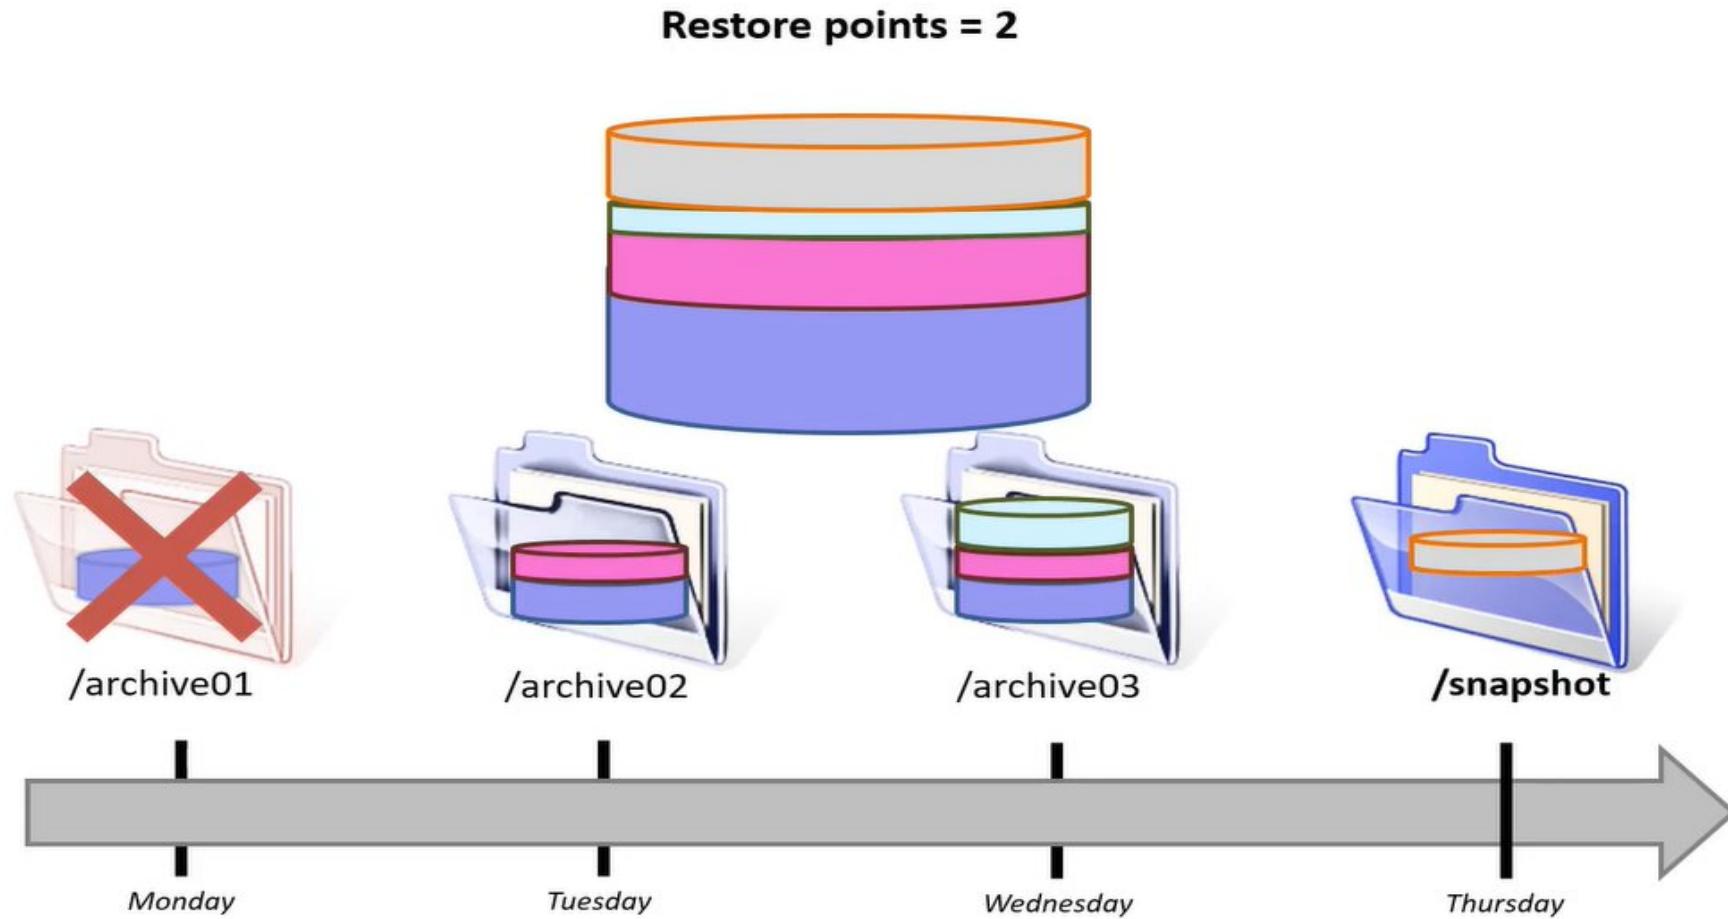
Restoring the Database from a Backup



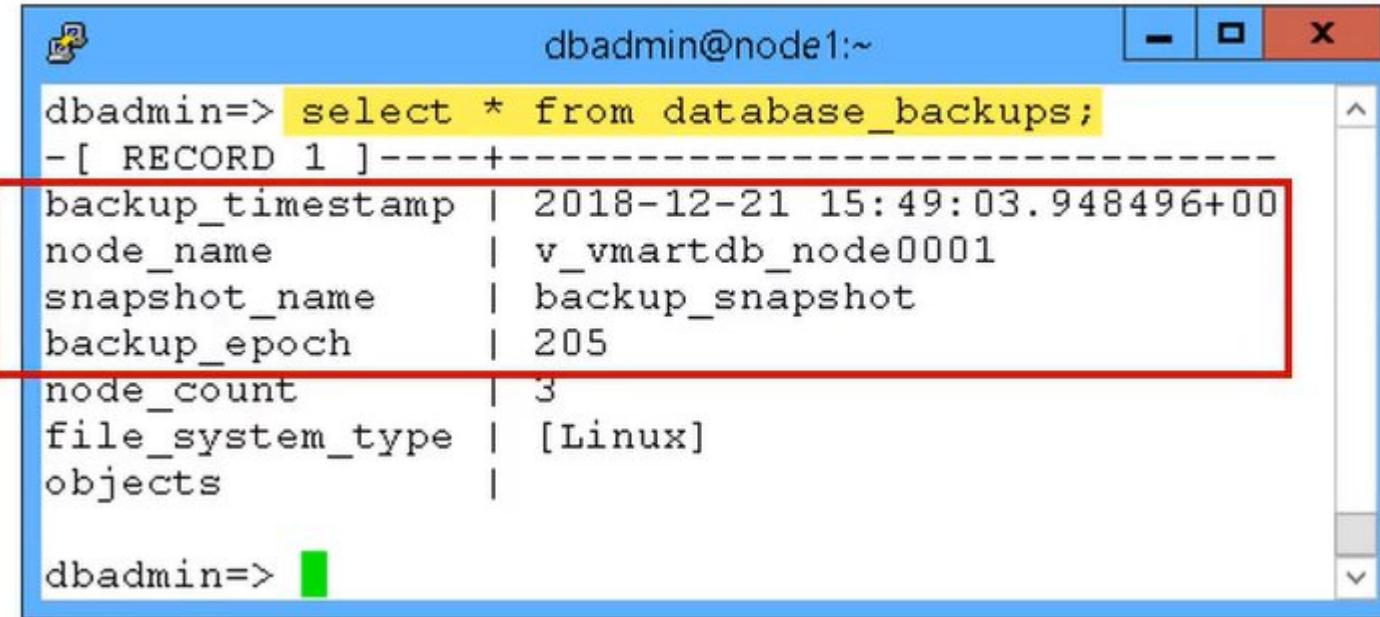
The screenshot shows a terminal window titled "dbadmin@node1:~". The command entered is "vbr --task restore --config-file ~/backups/backup_config.ini". The output indicates a full restore of the database VMartDB, involving three participating nodes: v_vmartdb_node0001, v_vmartdb_node0002, and v_vmartdb_node0003. The restore point used is backup_snapshot_20171101_204701. The process involves determining what data to restore, copying approximately 894912483 bytes of data (out of a total of 1601265689), syncing data from the backup to cluster nodes, and finally restoring the catalog. The restore is completed successfully.

```
[dbadmin@node1 ~]$ vbr --task restore --config-file ~/backups/backup_config.ini
Starting full restore of database VMartDB.
Participating nodes: v_vmartdb_node0001, v_vmartdb_node0002, v_vmartdb_node0003.
Restoring from restore point: backup_snapshot_20171101_204701
Determining what data to restore from backup.
[=====] 100%
Approximate bytes to copy: 894912483 of 1601265689 total.
Syncing data from backup to cluster nodes.
[=====] 100%
Restoring catalog.
Restore complete!
[dbadmin@node1 ~]$
```

Backup Restore points



Monitoring Backups



A screenshot of a terminal window titled "dbadmin@node1:~". The window contains the following command and its output:

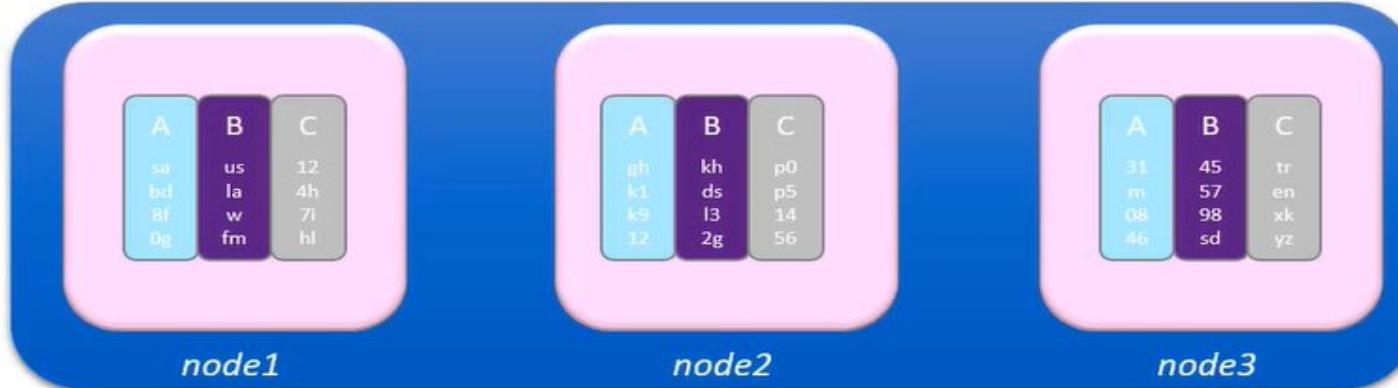
```
dbadmin=> select * from database_backups;
-[ RECORD 1 ]-----+
backup_timestamp | 2018-12-21 15:49:03.948496+00
node_name         | v_vmartdb_node0001
snapshot_name     | backup_snapshot
backup_epoch      | 205
node_count        | 3
file_system_type  | [Linux]
objects           |
```

The entire output table is highlighted with a red rectangular box.

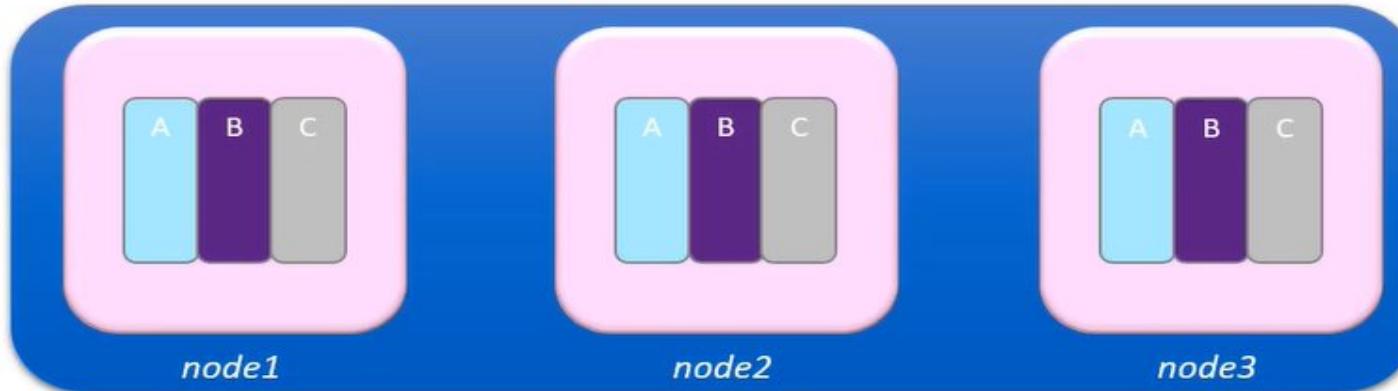
Restoring Objects

```
vbr --task restore --config-file filename.ini --restore-objects obj1..objN
```

Cluster1
original cluster
Status: UP



Cluster1
original cluster
Status: UP



Copy Vertica database

- Target cluster requirements:
 - Same number of nodes source cluster
 - Database with the same name as the database being copied
 - Same node names as the source cluster
 - Same database administrator account
- `vbr.py - -task copycluster - -config-file <configfile>`
- Ideal for creating a dormant Disaster Recovery site

Node Recovery (1 of 2)

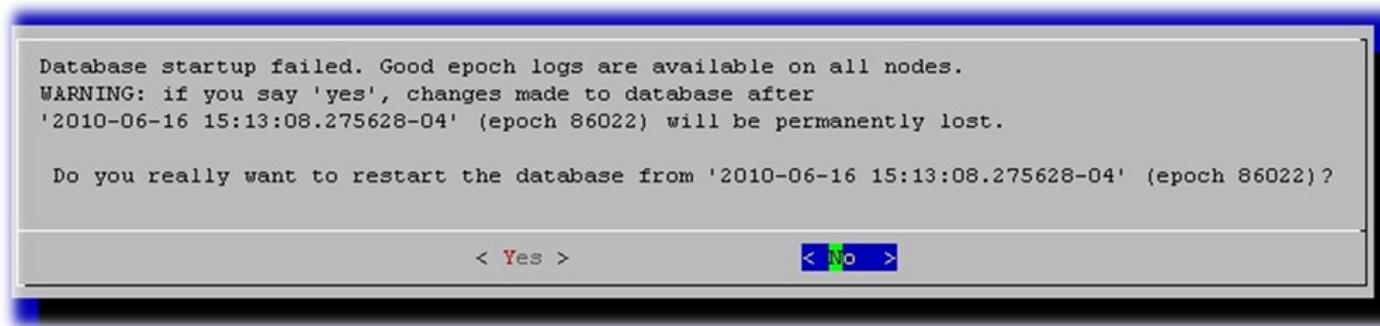
- A node can rebuild its data set from other nodes in the cluster if the cluster is K-safe
- Full recovery
 - Node rebuilds from scratch

Node Recovery (2 of 2)

- Incremental recovery
 - Node rebuilds from current persisted state
 - To speed up a full recovery, use a prior backup for the given node and perform incremental recovery
- RAID arrays (5,6,10) can be rebuilt without impact to other cluster nodes

Recover from Last Good Epoch (LGE)

If nodes contain persistent data from different epochs, the Last Good Epoch (LGE) on all nodes is used to determine recovery point



```
*** Restarting database db at epoch 86022 ***  
  
Node Status: v_db_node0001: (DOWN) v_db_node0002: (DOWN) v_db_node0003: (DOWN)  
  
Node Status: v_db_node0001: (INITIALIZING) v_db_node0002: (INITIALIZING) v_db_node0003: (INITIALIZING)  
  
Node Status: v_db_node0001: (RECOVERING) v_db_node0002: (RECOVERING) v_db_node0003: (RECOVERING)  
  
Node Status: v_db_node0001: (UP) v_db_node0002: (UP) v_db_node0003: (UP)
```

Monitoring Recovery

- Monitor disk space
 - df -h
 - SELECT * FROM v_monitor.disk_storage;
- Monitor recovery
 - tail -f <catalog-directory-path>/vertica.log

The background features a series of glowing blue light streaks and curves, resembling motion blur from a camera or light painting, set against a dark, solid black background.

opentext™

Thank you